

PROFESSIONAL SUMMARY

Data Scientist and ML Engineer with research and production experience across machine learning, LLM systems, and MLOps. Experienced deploying distributed ML pipelines and GPU-accelerated LLM serving on Kubernetes, building agentic AI applications, and developing novel ML methods published in peer-reviewed venues. Strong foundation in statistical modeling, deep learning, and time-series forecasting.

EDUCATION

- Master of Science in Data Science** Berlin, Germany
Berlin University of Applied Sciences (BHT) — Graduated With Distinction (1.06 German) Oct. 2024 – Sept. 2026
- Bachelor of Science in Computer Science; Minors in Math & Stats** Bellingham, WA
Western Washington University — GPA: 3.96 Sept. 2020 – June 2024

EXPERIENCE

- Data Scientist / ML Researcher** Berlin, Germany
Calgo Lab — Berlin University of Applied Sciences (BHT) — Studentische Hilfskraft Feb. 2025 – Sept. 2026
 - Engineered distributed ML experiment pipelines on a Kubernetes cluster with Helm; automated testing and linting via GitHub Actions CI/CD.
 - Deployed and maintained Dockerized LLM serving stack (vLLM on NVIDIA GPUs) and web-based research demos, including a demo used by Charité.
 - Architected **MechDetect**, a novel ML framework for classifying error mechanisms in tabular datasets; published at IEEE DSIS 2025.
 - Implemented spatiotemporal groundwater-level forecasting (master’s thesis) combining Chronos2, N-HiTS, and GPyTorch Gaussian Processes.
- Machine Learning Researcher** Bellingham, WA
Hutch Research — Western Washington University Mar. 2023 – June 2024
 - Built a PyTorch/TensorFlow probabilistic neural network pipeline to predict binary star system parameters; published in *The Astronomical Journal*, Vol. 171 (2026).
 - Optimized large astronomical dataset processing with NumPy; maintained code quality and reproducibility in a 4-person team.
- Statistics Researcher** Bellingham, WA
Western Washington University Oct. 2022 – June 2024
 - Developed a novel class of non-parametric two-sample tests in R; presented at Joint Mathematics Meetings 2024 and SIAM PNW 2023.

SELECTED PROJECTS

- LangGraph Email Agent** | vLLM, LangGraph, Kubernetes, Docker, Nvidia A100
 - Deployed an agentic LLM email-answering system on Kubernetes using vLLM for inference serving and LangGraph for multi-step orchestration with structured output.
- Tax Law RAG System** | BERT Fine-tuning, LangChain, FAISS, Kubernetes
 - Built a production RAG pipeline with fine-tuned BERT retrieval; benchmarked FAISS vs. Annoy vector databases on English and German corpora. Published the fine-tuned BERT to HuggingFace.
- Berlin PM2.5 Spatiotemporal Modeling** | Chronos2, Kubernetes, Docker, Streamlit
 - Automated daily ETL and inference pipeline on Kubernetes; spatiotemporal forecasting via Chronos2 foundation model and ordinary kriging.
- tab-err: Tabular Data Error Library** | Python, PyPI, CI/CD, GitHub Actions
 - Contributed to open-source library for realistic tabular error generation; developed high-level APIs and comprehensive unit tests.

PUBLICATIONS

- Anonymized but Useful Synthetic Tabular Health Data for AI based Fall Risk Assessment** | *Scientific Data* 2026
- Probabilistic Neural Networks for Eclipsing Binary Parameters** | *The Astronomical Journal* 2026
- MechDetect: Detecting Data-Dependent Errors** | IEEE DSIS 2025
- Towards Realistic Error Models for Tabular Data** | *ACM JDIQ* 2025

SKILLS

MLOps & Deployment: Kubernetes, Docker, Helm, vLLM, FastAPI, GitHub Actions, Prometheus, Grafana
LLM & Agentic AI: LangChain, LangGraph, FAISS, RAG, Prompt Engineering, Structured Outputs, Agentic Pipelines, Transformer Fine-tuning
ML & Data: PyTorch, TensorFlow, GPyTorch, Scikit-learn, Pandas, NumPy, PostgreSQL, MySQL, Time-Series, Spatial Data
Cloud & DevOps: AWS, GPU Workloads (NVIDIA), Containerization, Git, Linux
Languages: Python, SQL, R, Go, Bash, Java **Spoken:** English (Native), German (B1)